

Optional Lab: Cost Function for Logistic Regression

Goals

In this lab, you will:

- examine the implementation and utilize the cost function for logistic regression.

```
In [1]: import numpy as np
import matplotlib widget
import matplotlib.pyplot as plt
from lab_utils_common import plot_data, sigmoid, dlc
plt.style.use('./deeplearning.mplstyle')
```

Dataset


Let's start with the same dataset as was used in the decision boundary lab.

```
In [2]: X_train = np.array([[0.5, 1.5], [1, 1], [1.5, 0.5], [3, 0.5], [2, 2], [1, 2.5]]) # (m, n)
y_train = np.array([0, 0, 0, 1, 1, 1]) # (m,)
```

We will use a helper function to plot this data. The data points with label $y = 1$ are shown as red crosses, while the data points with label $y = 0$ are shown as blue circles.

```
In [3]: fig, ax = plt.subplots(1, 1, figsize=(4, 4))
plot_data(X_train, y_train, ax)

# Set both axes to be from 0-4
ax.axis([0, 4, 0, 3.5])
ax.set_ylabel('$x_1$', fontsize=12)
ax.set_xlabel('$x_0$', fontsize=12)
plt.show()
```



Cost function

In a previous lab, you developed the *logistic loss* function. Recall, loss is defined to apply to one example. Here you combine the losses to form the **cost**, which includes all the examples.

Recall that for logistic regression, the cost function is of the form

$$J(\mathbf{w}, b) = \frac{1}{m} \sum_{i=0}^{m-1} [\text{loss}(f_{\mathbf{w},b}(\mathbf{x}^{(i)}), y^{(i)})] \quad (1)$$

where

- $\text{loss}(f_{\mathbf{w},b}(\mathbf{x}^{(i)}), y^{(i)})$ is the cost for a single data point, which is:

$$\text{loss}(f_{\mathbf{w},b}(\mathbf{x}^{(i)}), y^{(i)}) = -y^{(i)} \log(f_{\mathbf{w},b}(\mathbf{x}^{(i)})) - (1 - y^{(i)}) \log(1 - f_{\mathbf{w},b}(\mathbf{x}^{(i)})) \quad (2)$$

- where m is the number of training examples in the data set and:

$$f_{\mathbf{w},b}(\mathbf{x}^{(i)}) = g(z^{(i)}) \quad (3)$$

$$z^{(i)} = \mathbf{w} \cdot \mathbf{x}^{(i)} + b \quad (4)$$

$$g(z^{(i)}) = \frac{1}{1 + e^{-z^{(i)}}} \quad (5)$$

Code Description

The algorithm for `compute_cost_logistic` loops over all the examples calculating the loss for each example and accumulating the total.

Note that the variables X and y are not scalar values but matrices of shape (m, n) and $(m,)$ respectively, where n is the number of features and m is the number of training examples.

```
In [4]: def compute_cost_logistic(X, y, w, b):  
    """  
    Computes cost  
  
    Args:  
        X (ndarray (m,n)): Data, m examples with n features  
        y (ndarray (m,)) : target values  
        w (ndarray (n,)) : model parameters  
        b (scalar)       : model parameter  
  
    Returns:  
        cost (scalar): cost  
    """  
  
    m = X.shape[0]  
    cost = 0.0  
    for i in range(m):  
        z_i = np.dot(X[i], w) + b  
        f_wb_i = sigmoid(z_i)  
        cost += -y[i]*np.log(f_wb_i) - (1-y[i])*np.log(1-f_wb_i)  
  
    cost = cost / m  
    return cost
```

Check the implementation of the cost function using the cell below.

```
In [5]: w_tmp = np.array([1,1])  
        b_tmp = -3  
        print(compute_cost_logistic(X_train, y_train, w_tmp, b_tmp))  
  
0.36686678640551745
```

Expected output: 0.3668667864055175

Example


Now, let's see what the cost function output is for a different value of w .

- In a previous lab, you plotted the decision boundary for $b = -3, w_0 = 1, w_1 = 1$. That is, you had $b = -3, w = \text{np.array}([1,1])$.
- Let's say you want to see if $b = -4, w_0 = 1, w_1 = 1$, or $b = -4, w = \text{np.array}([1,1])$ provides a better model.

Let's first plot the decision boundary for these two different b values to see which one fits the data better.

- For $b = -3, w_0 = 1, w_1 = 1$, we'll plot $-3 + x_0 + x_1 = 0$ (shown in blue)
- For $b = -4, w_0 = 1, w_1 = 1$, we'll plot $-4 + x_0 + x_1 = 0$ (shown in magenta)

```
In [6]: import matplotlib.pyplot as plt  
  
# Choose values between 0 and 6  
x0 = np.arange(0,6)  
  
# Plot the two decision boundaries  
x1 = 3 - x0  
x1_other = 4 - x0  
  
fig, ax = plt.subplots(1, 1, figsize=(4,4))  
# Plot the decision boundary  
ax.plot(x0, x1, c='dblue', label="$b=-3$")  
ax.plot(x0, x1_other, c='dlmagenta', label="$b=-4$")  
ax.axis([0, 4, 0, 4])  
  
# Plot the original data  
plot_data(X_train, y_train, ax)  
ax.axis([0, 4, 0, 4])  
ax.set_ylabel('$x_1$', fontsize=12)  
ax.set_xlabel('$x_0$', fontsize=12)  
plt.legend(loc="upper right")  
plt.title("Decision Boundary")  
plt.show()
```



You can see from this plot that $b = -4, w = \text{np.array}([1,1])$ is a worse model for the training data. Let's see if the cost function implementation reflects this.

```
In [7]: w_array1 = np.array([1,1])
        b_1 = -3
        w_array2 = np.array([1,1])
        b_2 = -4

        print("Cost for b = -3 : ", compute_cost_logistic(X_train, y_train, w_array1, b_1))
        print("Cost for b = -4 : ", compute_cost_logistic(X_train, y_train, w_array2, b_2))

Cost for b = -3 : 0.36686678640551745
Cost for b = -4 : 0.5036808636748461
```

Expected output

Cost for b = -3 : 0.3668667864055175

Cost for b = -4 : 0.5036808636748461

You can see the cost function behaves as expected and the cost for `b = -4, w = np.array([1,1])` is indeed higher than the cost for `b = -3, w = np.array([1,1])`

Congratulations!

In this lab you examined and utilized the cost function for logistic regression.

In []: